

# Introduction to OpenTelemetry with Python

# Jessica Garson

@JessicaGarson

@JessicaGarson@macaw.social

Senior Developer Advocate



<https://github.com/JessicaGarson/Introduction-to-OpenTelemetry-with-Python>



# Agenda

- Observability basics
- What is OpenTelemetry and why use it
- Django and Flask instrumentation
- The state of OpenTelemetry
- Conclusion and next steps

I had attempted to find a short primer on starting with OpenTelemetry with Python but couldn't find much more than a [demo application](#).

<https://opentelemetry-python.readthedocs.io/en/latest/examples/django/README.html>  
<https://opentelemetry.io/docs/languages/python/getting-started/>





## ChatGPT

Sure! Imagine you have a big box of crayons (these are like different parts of a computer program) and you want to make a beautiful picture (which is like running the program to do something cool). But sometimes, while you're drawing, a crayon might break or not color the way it should. This can make your picture not turn out as nice as you wanted.

OpenTelemetry is like having a magical friend who watches you while you draw. This friend takes notes on which crayons you use, how they perform, and if they break. Then, if something goes wrong, your magical friend can tell you exactly which crayon broke and why, so you can fix it and make sure your picture turns out great next time!

So, OpenTelemetry helps people who make computer programs make sure their programs run smoothly and fix any problems that happen.



# Observability fundamentals

Observability helps you determine what's happening inside your system by analyzing the data it outputs



# What isn't operating as expected

# Beyond detecting problems

# What is an agent?

An observability agent is a software component or tool that collects, processes, and sends telemetry data from a system or application to an observability platform for analysis.

# Logs

Logs are records of events in a system, documenting operations, errors, and activities to aid in troubleshooting, monitoring, and compliance.

# Metrics

Metrics are quantitative measurements that track the performance and health of a system

# Traces

Traces track the path and interactions of a request through a system.

# What is OpenTelemetry (OTel)?

**OpenTelemetry (OTel) is an observability framework that is open source and vendor neutral, designed to work with any backend system.**

It provides standardized APIs, libraries, and tools to collect telemetry data such as metrics, logs, and traces.

OpenTelemetry enables comprehensive monitoring  
across services and platforms.

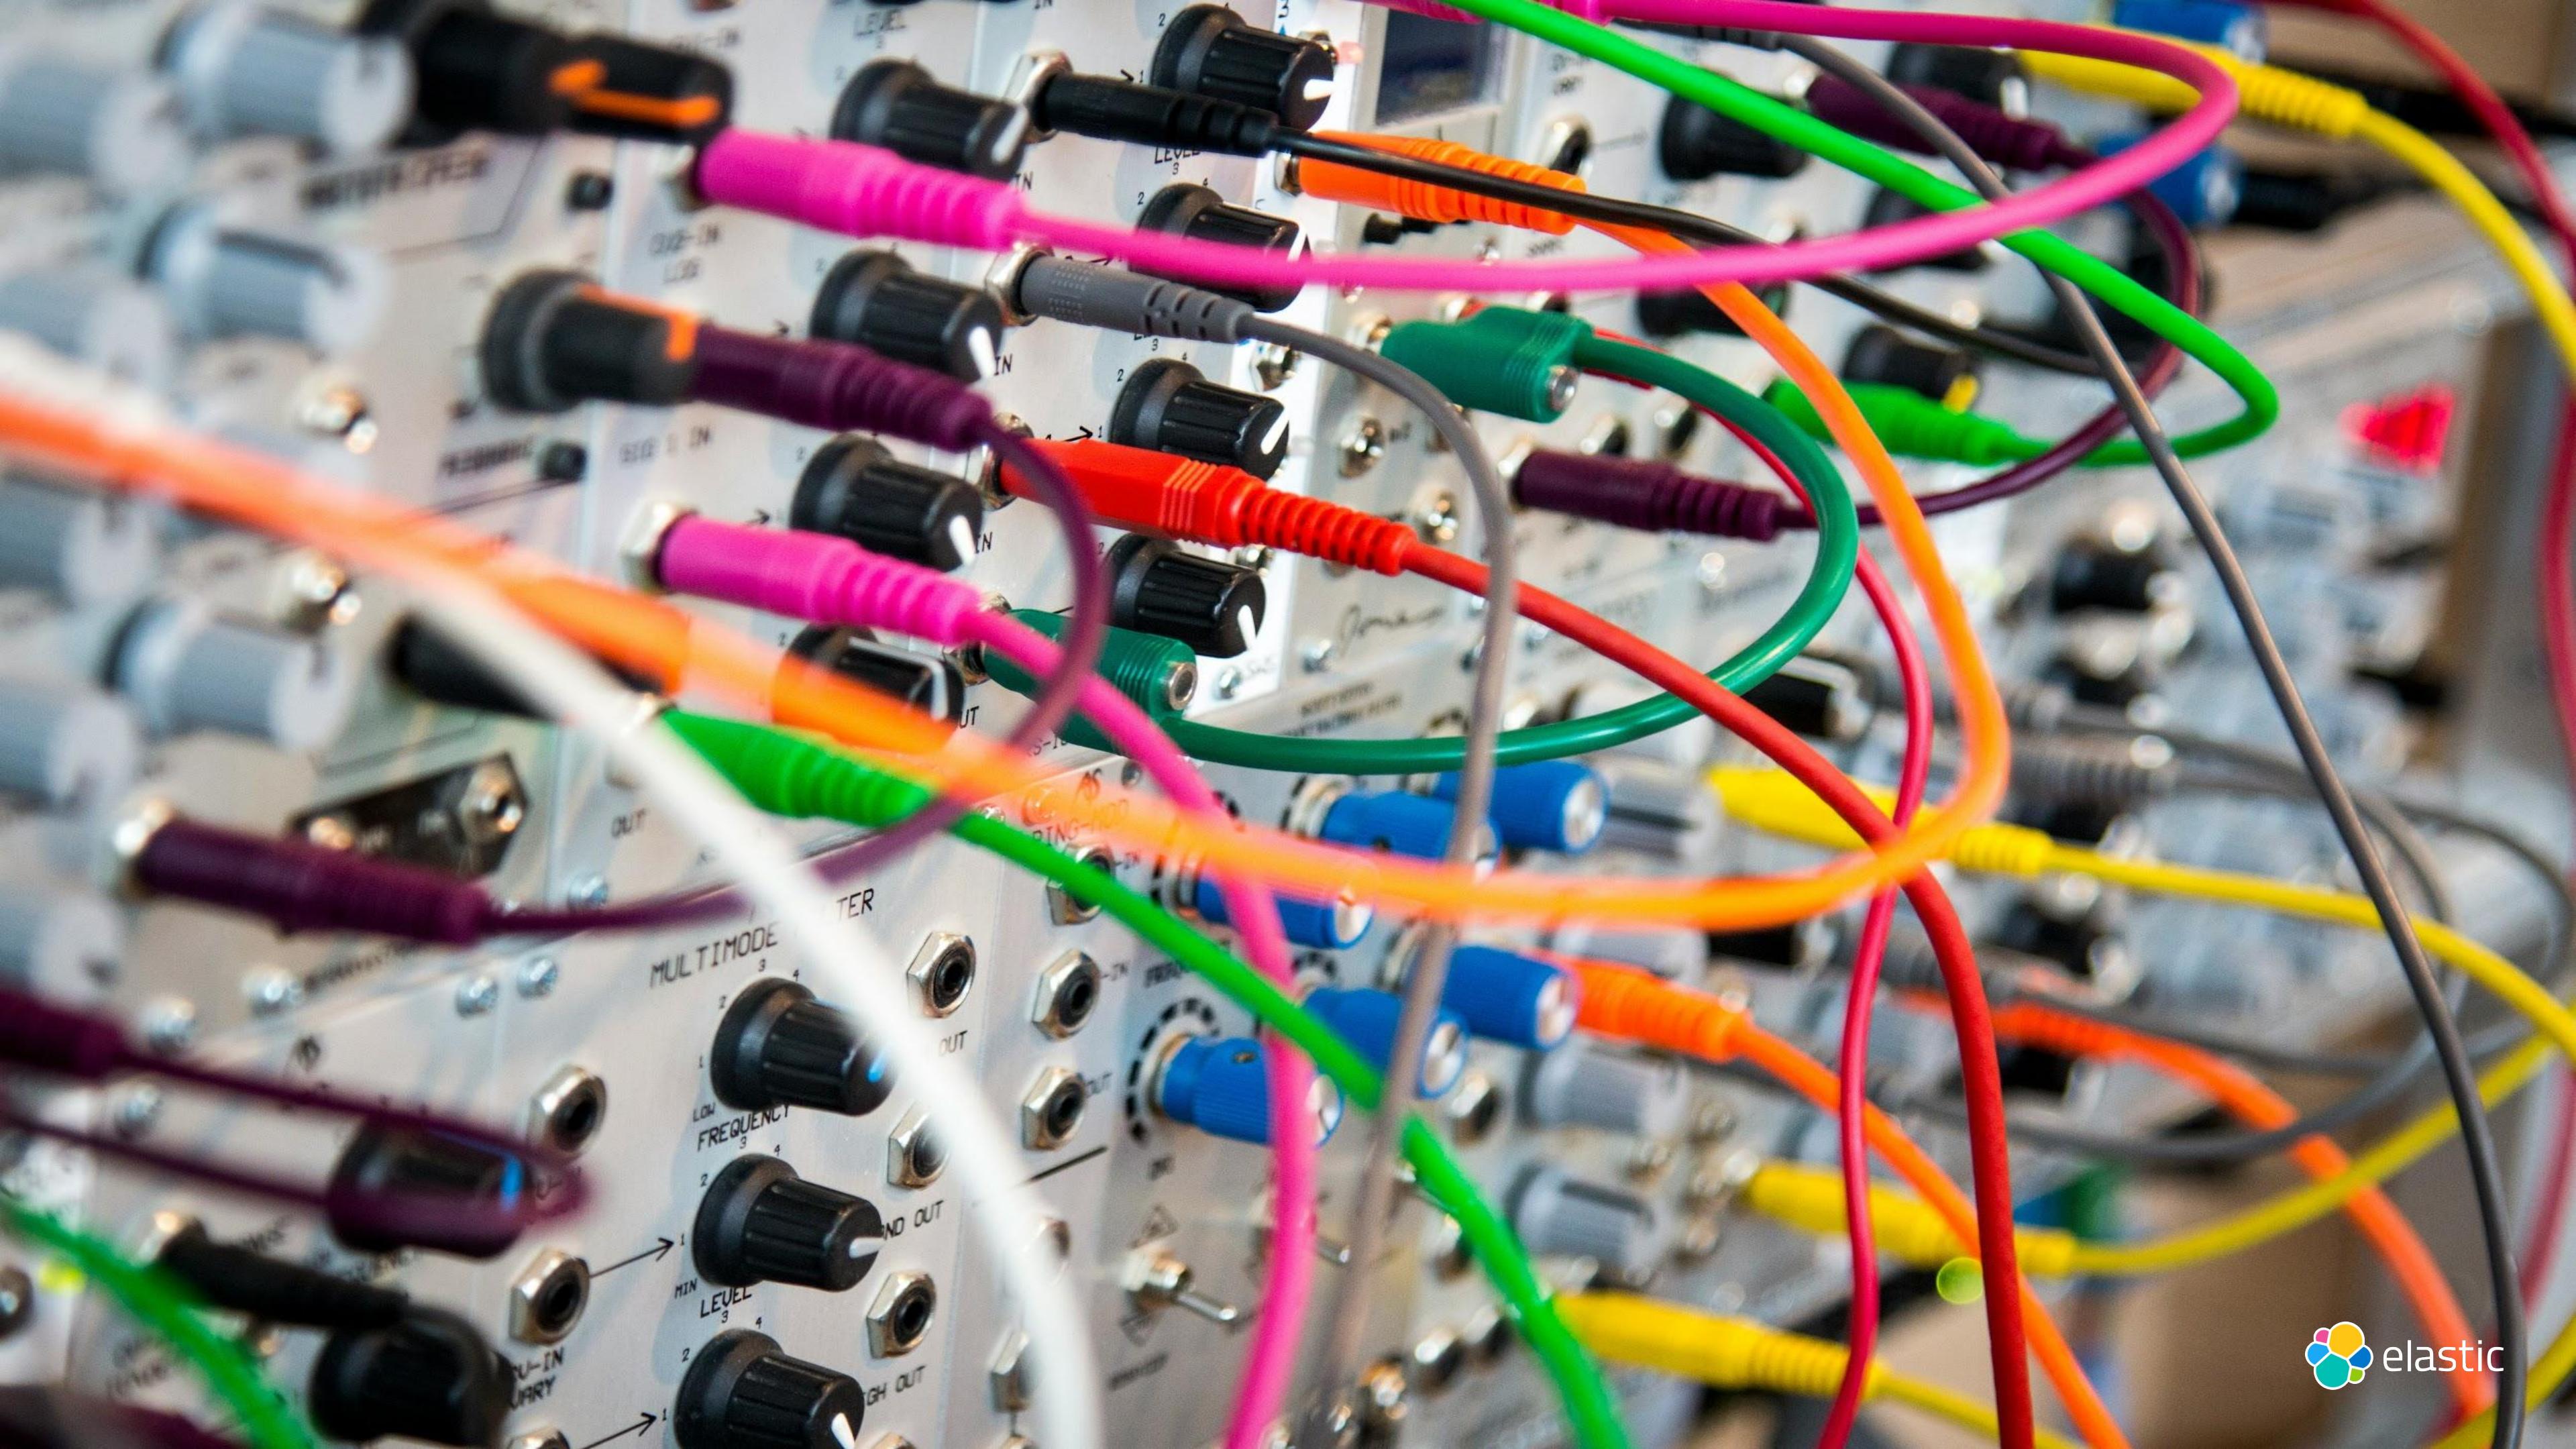
# Why?

A photograph of a stack of logs in a forest. The logs are cut trees, showing their circular cross-sections. They are piled up in a somewhat haphazard manner, with some logs leaning against each other. The background is filled with dense, dark forest trees. The lighting suggests it might be late afternoon or early morning, with dappled sunlight filtering through the canopy.

Logs alone are simply not  
enough sometimes

# Um, what?

>	*	2024-05-10	08:00:06.117	55	54440319	...	http://api.nasa.gov/neo/rest/v1/neo/sentry/544...	
>	*	2024-05-10	08:00:06.117	56	54440642	...		Nan
▼	*	2024-05-10	08:00:06.117	57	54440644	...	http://api.nasa.gov/neo/rest/v1/neo/sentry/544...	
<div><span>Copy</span> <span>Similar entries</span> <span>Expand nested fields</span> <span>Hide log summary</span></div>								
▼	{							
	insertId:	"663e0c460001c9a60e5e65c3"						
▶	labels:	{2}						
	logName:	"projects/ncav-293119/logs/run.googleapis.com%2Fstdout"						
	receiveTimestamp:	"2024-05-10T12:00:06.338017672Z"						
▶	resource:	{2}						
	textPayload:	"57 54440644 ... http://api.nasa.gov/neo/rest/v1/neo/sentry/544..."						
	timestamp:	"2024-05-10T12:00:06.117158Z"						
	}							
>	*	2024-05-10	08:00:06.117	[43 rows x 18 columns]				



# Future proof



# Spans within traces? What?!?!

Spans are the building blocks of traces.  
They include the following information:

- Name
- Parent span ID (empty for root spans)

```
{  
  "name": "add_item_view_span",  
  "context": {  
    "trace_id": "0x81364006fe668cb26a54e20918be9620",  
    "span_id": "0x9a048a860bb7c9ac",  
    "trace_state": "[]"  
  },  
  "kind": "SpanKind.INTERNAL",  
  "parent_id": null,  
  "start_time": "2024-08-15T20:55:42.634538Z",  
  "end_time": "2024-08-15T20:55:42.637830Z",  
  "status": {  
    "status_code": "UNSET"  
  },  
  "attributes": {},  
  "events": [],  
  "links": [],  
  "resource": {  
    "attributes": {  
      "service.name": "to-do-list-app-manual",  
      "service.version": "1.0.0"  
    },  
    "schema_url": ""  
  }  
}
```

OHHH!

A span  
within a  
trace!

Instrumenting refers to the process of adding observability features to your application to collect telemetry data, such as traces, metrics, and logs.

# Django Example

# The classic Django example

**To Do List**

- New task! - Cool! A new task! Delete
- Another task! - Yay! Delete

**Add New Item**

# How do you get that data to show in your console?



```
pip install django django-environ elastic-opentelemetry opentelemetry-instrumentation-django  
opentelemetry-bootstrap --action=install
```



# settings.py



```
env = environ.Env( )

environ.Env.read_env(os.path.join(BASE_DIR, '.env'))

OTEL_EXPORTER_OTLP_HEADERS = env('OTEL_EXPORTER_OTLP_HEADERS')
OTEL_EXPORTER_OTLP_ENDPOINT = env('OTEL_EXPORTER_OTLP_ENDPOINT')
```

# .env example



```
OTEL_EXPORTER_OTLP_HEADERS="Authorization=ApiKey%20yourapikey"  
OTEL_EXPORTER_OTLP_ENDPOINT="https://your/host/endpoint"
```



```
import os
import sys
from opentelemetry import trace
from opentelemetry.sdk.trace import TracerProvider
from opentelemetry.sdk.trace.export import BatchSpanProcessor, ConsoleSpanExporter
from opentelemetry.sdk import resources

def main():
    """Run administrative tasks."""
    os.environ.setdefault("DJANGO_SETTINGS_MODULE", "todolist_project.settings")

    resource = resources.Resource(attributes={
        resources.SERVICE_NAME: "your-service-name",
        resources.SERVICE_VERSION: "1.0.0"
    })

    trace_provider = TracerProvider(resource=resource)
    trace.set_tracer_provider(trace_provider)

    console_exporter = ConsoleSpanExporter()

    span_processor = BatchSpanProcessor(console_exporter)
    trace.get_tracer_provider().add_span_processor(span_processor)
```

# manage.py



# Automatic Instrumentation

The process by which an agent modifies the bytecode of your application's classes, often to insert monitoring code

# manage.py

```
def main():
    """Run administrative tasks."""
    os.environ.setdefault("DJANGO_SETTINGS_MODULE", "todolist_project.settings")
    DjangoInstrumentor().instrument()

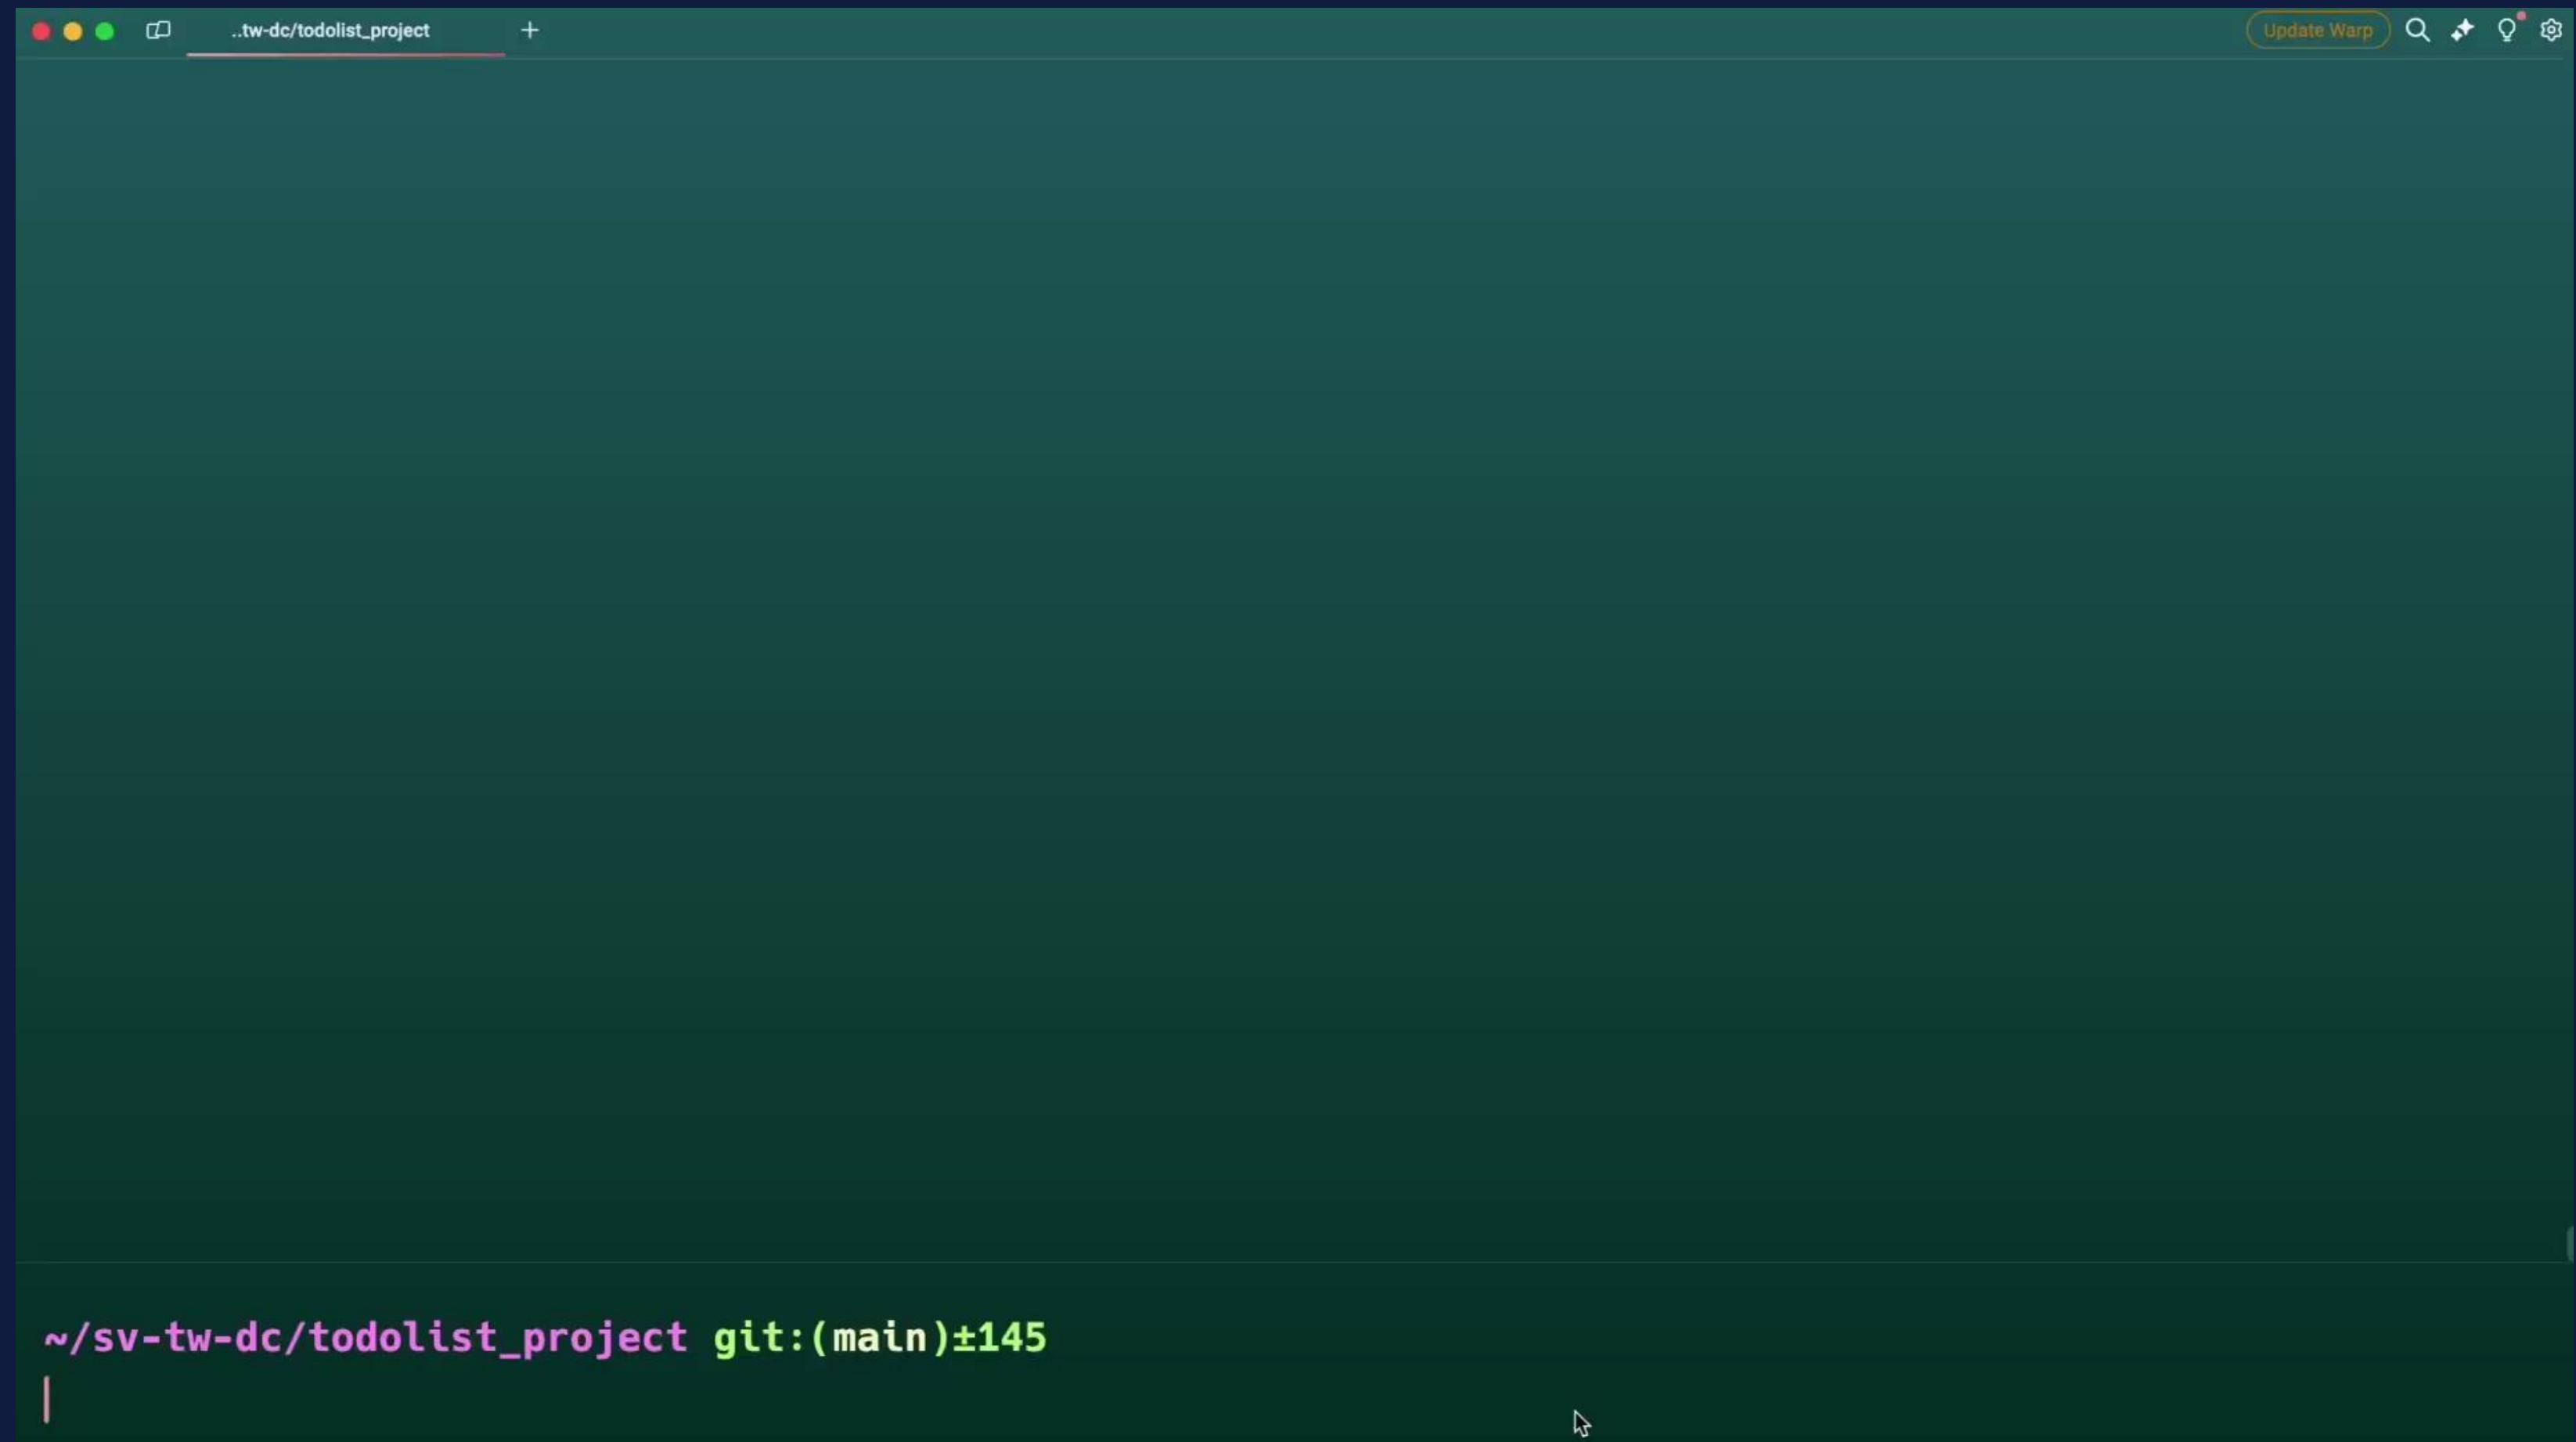
    # Set up resource attributes for the service
    resource = resources.Resource(attributes={
        resources.SERVICE_NAME: "to-do-list-app",
        resources.SERVICE_VERSION: "1.0.0"
    })

    trace_provider = TracerProvider(resource=resource)
    trace.set_tracer_provider(trace_provider)

    otlp_exporter = OTLPSpanExporter()

    # Set up the BatchSpanProcessor to export traces
    span_processor = BatchSpanProcessor(otlp_exporter)
    trace.get_tracer_provider().add_span_processor(span_processor)
    turn.go(f, seed, [])
}
```





A screenshot of a terminal window with a dark background. The title bar at the top shows the path `..tw-dc/todolist_project`. The bottom status bar displays the current directory as `~/sv-tw-dc/todolist_project git:(main)±145`. The main body of the window is entirely blank, indicating no code or output has been entered.



# Manual instrumentation

Manual instrumentation requires incorporating particular code segments into your application to collect and transmit telemetry data.

**So when would you want to use manual instrumentation?**

# Step 1: Delete this line from your manage.py file



```
# DjangoInstrumentor().instrument()
```

## Step 2: Update your views.py

```
tracer = trace.get_tracer(__name__)

meter = get_meter(__name__)

view_counter = meter.create_counter(
    "view_requests",
    description="Counts the number of requests to views",
)

view_duration_histogram = meter.create_histogram(
    "view_duration",
    description="Measures the duration of view execution",
)
```



```
def index(request):
    start_time = time()

    with tracer.start_as_current_span("index_view_span") as span:
        items = ToDoItem.objects.all()
        span.set_attribute("todo.item_count", items.count())
        response = render(request, 'todo/index.html', {'items': items})

    view_counter.add(1, {"view_name": "index"})
    view_duration_histogram.record(time() - start_time, {"view_name": "index"})

    return response
```



```
def add_item(request):
    start_time = time()

    with tracer.start_as_current_span("add_item_view_span") as span:
        if request.method == 'POST':
            form = ToDoForm(request.POST)
            if form.is_valid():
                form.save()
                span.add_event("New item added")
                response = redirect('index')
            else:
                response = render(request, 'todo/add_item.html', {'form': form})
        else:
            form = ToDoForm()
            response = render(request, 'todo/add_item.html', {'form': form})

    view_counter.add(1, {"view_name": "add_item"})
    view_duration_histogram.record(time() - start_time, {"view_name": "add_item"})

    return response
```





```
def delete_item(request, item_id):
    start_time = time()

    with tracer.start_as_current_span("delete_item_view_span") as span:
        item = get_object_or_404(TodoItem, id=item_id)
        span.set_attribute("todo.item_id", item_id)
        if request.method == 'POST':
            item.delete()
            span.add_event("Item deleted")
            response = redirect('index')
        else:
            response = render(request, 'todo/delete_item.html', {'item': item})

    view_counter.add(1, {"view_name": "delete_item"})
    view_duration_histogram.record(time() - start_time, {"view_name": "delete_item"})

return response
```



# Step 3: Update your models.py

```
from django.db import models
from opentelemetry import trace

tracer = trace.get_tracer(__name__)

class ToDoItem(models.Model):
    title = models.CharField(max_length=100)
    description = models.TextField(blank=True)
    created_at = models.DateTimeField(auto_now_add=True)

    def __str__(self):
        return self.title

    def save(self, *args, **kwargs):
        with tracer.start_as_current_span("save_todo_item_span") as span:
            span.set_attribute("todo.title", self.title)
            if self.pk:
                span.add_event("Updating ToDoItem")
            else:
                span.add_event("Creating new ToDoItem")

        super(ToDoItem, self).save(*args, **kwargs)
```



```
python manage.py runserver
```

```
Watching for file changes with StatReloader  
Performing system checks...
```

```
System check identified no issues (0 silenced).
```

```
October 23, 2024 - 23:36:52
```

```
Django version 5.1, using settings 'todolist_project.settings'
```

```
Starting development server at http://127.0.0.1:8000/
```

```
Quit the server with CONTROL-C.
```

```
[23/Oct/2024 23:36:59] "GET / HTTP/1.1" 200 1695
```

# Flask

# No code instrumentation for automatic instrumentation



pip install elastic-opentelemetry



```
opentelemetry-bootstrap --action=install
```



```
export OTEL_RESOURCE_ATTRIBUTES="service.name=flask-todo-app"
export OTEL_EXPORTER_OTLP_HEADERS="Authorization=<your_authorization_header_value>"
export OTEL_EXPORTER_OTLP_ENDPOINT="<your_elastic_cloud_url>"
```





opentelemetry-instrument flask run

# To-Do List

Add new task

ADD TASK

New task!

Delete

# Manual instrumentation in Flask

~/python-otel/manual\_flask git:(main)±154

# OpenTelemetry Collector

A solution for receiving, processing, and exporting telemetry data. It eliminates the need to manage and maintain multiple agents or collectors.

<https://opentelemetry.io/docs/demo>



elastic / opentelemetry-demo

 Type / to search

Code

Issues 4

Pull requests 26

Actions

Projects

Security

Insights



opentelemetry-demo

Public

forked from [open-telemetry/opentelemetry-demo](#)

Watch 5

Fork 28

Star 21

main

24 Branches



4 Tags

Go to file



Add file

Code

This branch is 94 commits ahead of [open-telemetry/opentelemetry-demo:main](#).

#80

	github-actions Merge remote-tracking branch 'upstream/main' ✓	f3c606c · 5 hours ago	988 Commits
	.github Merge remote-tracking branch 'upstream/main'	5 hours ago	
	internal/tools Update wiki url (open-telemetry#1346)	8 months ago	
	kubernetes remove configMap k8s Helm dependency (#77)	2 weeks ago	
	pb [ffs] - Allow setting initial feature flag values (open-telemetry#1713)	8 months ago	
	src Merge remote-tracking branch 'upstream/main'	7 hours ago	
	test Bump dependencies (open-telemetry#1713)	last month	
	.dockerignore Add license check (open-telemetry#825)	last year	
	.env Bump dependencies (open-telemetry#1713)	last month	
	.env.override remove configMap k8s Helm dependency (#77)	2 weeks ago	

## About

OpenTelemetry Community Demo Application

Readme

Apache-2.0 license

Security policy

Activity

Custom properties

21 stars

5 watching

28 forks

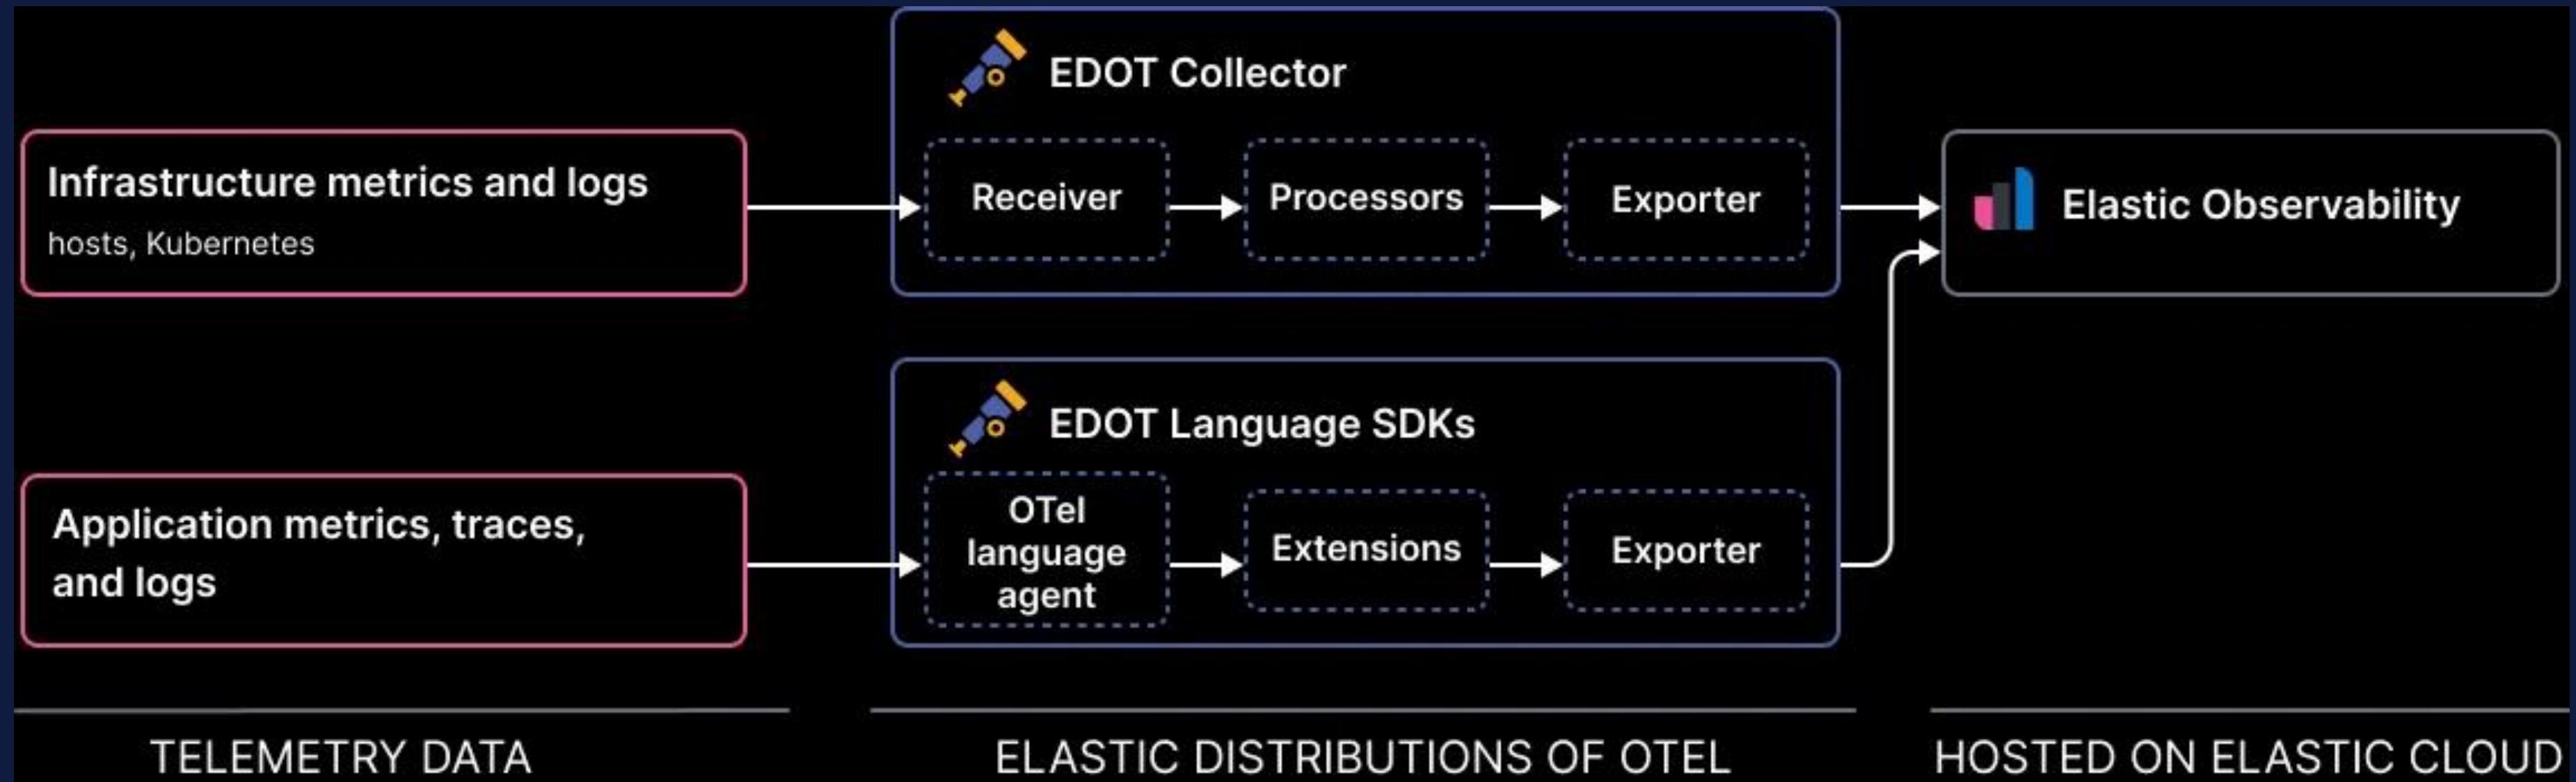
Report repository

## Releases 4

1.11.4 Latest

3 weeks ago

+ 3 releases



# State of adoption

WORK in PROGRESS

<https://opentelemetry.io/status/>

# Closing thoughts

OpenTelemetry (OTel) is an open-source,  
vendor-neutral observability framework  
designed to integrate with any backend  
system.

It's very easy to start instrumenting your code.

It is highly configurable and extensible

It scales well from large applications to  
small applications

**Let me know if this talk inspires you to build  
anything. I'm @JessicaGarson on most  
platforms.**

# Thank you!